

Pool Game: Concurrent Implementation Report

Assignment 01 - PCD Course

Academic Year 2025-2026

Abstract

This report presents a concurrent implementation of the Pool Game, a 2D board game where two players (human and bot) control balls to push smaller balls into holes. The document provides a detailed analysis of concurrency challenges, design architecture, behavioral specifications using Petri Nets, performance evaluation, and formal verification using JPF (Java Path Finder). The concurrent implementation demonstrates significant improvements over a sequential baseline while effectively exploiting available processor cores through multithreading and task-based approaches.

Contents

1	Problem Analysis	3
1.1	Problem Description	3
1.2	Multi-Threaded Related Challenges	3
1.2.1	Collision Detection Complexity	3
1.2.2	Memory Bandwidth Saturation	3
1.2.3	Synchronization Requirements	4
1.3	Task-based Related Challenges	4
1.3.1	Computational Complexity	5
1.3.2	Temporal Constraints	6
1.3.3	Synchronization Requirements	6
2	Design and Architecture	7
2.1	Architectural Overview for Multithreading	7
2.1.1	Thread Model	8
2.1.2	Parallel Physics Pipeline	9
2.1.3	Synchronization Mechanisms	11
2.1.4	MVC and Multithreading Interaction	13
2.2	Architectural Overview for Task-Based Approach	14
2.2.1	Component Descriptions	15
2.3	Synchronization Strategy	19
2.3.1	Barrier Synchronization via <code>invokeAll()</code>	19
2.3.2	Fine-Grained Ball Locking in Collision Detection	19
2.3.3	<code>ConcurrentHashMap</code> for Shared State	20
2.3.4	Exception Handling and Reliability	20
2.3.5	Memory Visibility and Happens-Before Guarantees	21

3	Behavior Specification: Petri Nets	21
3.1	Multi-Threaded CyclicBarrier Net	21
3.2	Multi-Threaded Tick Flow	22
3.3	High-Level Game Flow Net Task-based	23
4	Performance Evaluation	24
4.1	Experimental Setup	24
4.1.1	Test Configurations	24
4.1.2	Performance Metrics	25
4.2	Expected Performance Improvements	25
4.2.1	Sequential Baseline Issues	25
4.2.2	Concurrent Improvements	25
4.2.3	Task-based Improvements	26
5	Formal Verification with JPF	26
5.1	Java Path Finder Integration	26
5.1.1	JPF Configuration	26
5.2	Safety Properties Verified	26
5.2.1	1. Mutual Exclusion (No Data Races)	26
5.2.2	2. Deadlock Freedom	27
5.2.3	3. Liveness (Game Termination)	27
5.3	Race Condition Detection	27
5.3.1	Synchronized Field Accesses	27
5.3.2	Temporal Assertions	28
5.4	Model Checking Results	28
5.4.1	Test Configuration: MainJPF	28

1 Problem Analysis

1.1 Problem Description

The Pool Game is a two-player billiard-like game on a 2D board with the following characteristics:

- **Dynamic Objects:** Multiple balls (typically ranging from 2 to 4500) moving and interacting on the board
- **Physics Simulation:** Elastic collisions, friction forces, and momentum transfer between bodies
- **Players:** Two players (human controlled via keyboard, bot controlled autonomously) with scoring systems
- **Game Objective:** Insert small balls into holes at board corners; winner is determined by score or by eliminating opponent's ball

1.2 Multi-Threaded Related Challenges

Introducing parallelism into the physics pipeline exposed several concurrency challenges that influenced the overall architecture of the system. This section analyses each challenge and motivates the design choices.

1.2.1 Collision Detection Complexity

The most computationally demanding operation in the simulation is ball-ball collision detection. A naive approach tests every ball against every other ball each tick, yielding a time complexity of $O(n^2)$ where n is the number of small balls on the field. While parallelising this work across N workers reduces the per-thread cost to $O(n^2/N)$, the quadratic growth dominates: with a large number of balls, even distributing the work across all available cores is insufficient to maintain interactive framerates, because the total number of pair checks grows much faster than the number of available threads.

The underlying observation is that the vast majority of these checks are geometrically impossible: two balls of radius r can only collide if their centres are within distance $2r$ of each other. Checking distant pairs is wasted work. A spatial hashing grid — partitioning the field into cells of side $2r$ and testing each ball only against the balls in its 3×3 cell neighbourhood — reduces the average complexity to $O(n)$, since the number of neighbours per cell is bounded by a constant determined by ball density and radius. This algorithmic improvement is a prerequisite for parallelism to be effective at scale: no amount of additional threads can compensate for a fundamentally super-linear algorithm when n is large.

1.2.2 Memory Bandwidth Saturation

Distributing ball updates across multiple worker threads introduces the risk of memory bandwidth saturation and cache thrashing. If balls are assigned to workers randomly, the working set of each thread is scattered across the heap, causing frequent cache misses as each core loads and evicts cache lines that other cores also need. In the worst case, the

overhead of cache invalidation traffic between cores can exceed the computational benefit of parallelism.

A related issue arises from shared mutable state: the `lastTouchedBy` map, which records the last player to touch each ball for scoring purposes, is accessed by both the worker threads (to reset entries after ball-ball collisions) and the update thread (to update entries after player-ball collisions). Protecting this map with a single global lock on `Board` would serialise all workers on every collision, effectively eliminating concurrency for any tick where collisions are frequent. Furthermore, acquiring `synchronized(board)` while already holding `synchronized(ball)` introduces a lock ordering that can interleave incorrectly with other synchronized methods on `Board`, creating a potential deadlock.

Both issues are addressed by two targeted measures: assigning each worker a *contiguous* partition of the ball list (improving spatial locality and cache reuse), and replacing the `HashMap` with a `ConcurrentHashMap` (enabling concurrent updates without a global lock and eliminating the nested lock dependency).

1.2.3 Synchronization Requirements

Parallelising the physics pipeline requires coordinating multiple threads across several phases per tick, each with fundamentally different synchronization semantics. Three distinct requirements were identified.

Phase boundary enforcement. Position updates must be fully completed by all workers before any collision resolution begins. This is an *all-to-all* dependency: every thread must complete its current phase before any thread may advance to the next. The natural primitive is a `CyclicBarrier` — a reusable synchronization point that blocks all $N + 1$ parties (workers plus the `UpdateThread`) until the last one arrives, then releases all of them simultaneously. Its cyclic nature makes it well suited to the repeating tick structure without requiring reallocation.

Deadlock-free per-object locking. Concurrent collision resolution requires each worker to acquire locks on two `Ball` objects simultaneously. If different workers acquire these locks in inconsistent orders, circular waits can form — the classical lock inversion problem. The solution is a canonical lock-ordering protocol: locks are always acquired in strictly ascending index order (the position of each ball in the shared snapshot), which is a strict total order that prevents any circular dependency by construction.

Asymmetric synchronization for grid construction. The spatial grid is rebuilt each tick by a single designated thread (Worker-0), while all other threads must wait for it to be ready before entering the collision phase. This is a *one-to-all* dependency — one producer unblocks N consumers — which is semantically distinct from the all-to-all synchronization of a `CyclicBarrier`. Using a barrier here would be technically correct but conceptually misleading, as it implies parallel work where none exists. The appropriate primitive is a custom `ReusableLatch`: Worker-0 calls `open()` after `grid.build()`, all other threads call `await()` and block until the latch opens. The latch is reset by the `GameController` at the beginning of each tick, before any worker is activated, making it safely reusable across ticks without reallocation.

1.3 Task-based Related Challenges

The task-based concurrent implementation using the Executor Framework introduces a different set of challenges compared to the thread-per-entity model. The Pool imple-

mentation uses work-stealing task decomposition via `WorkerPool` to parallelize physics updates and collision detection.

1.3.1 Computational Complexity

Task-based parallelism in the Pool implementation must address several complexity-related challenges:

- **Task Granularity Trade-off:**

- Fine-grained tasks (small chunk size) improve load balancing but increase task creation and scheduling overhead
- In `WorkerPool`, each task updates balls in range $[i, i + \text{chunkSize})$, where chunk size is computed as $\text{chunkSize} = \max(1, n/k)$ with $k = 2 \cdot \text{numProcessors}$

- **Collision Detection Complexity:**

- `CollisionTask` must verify all pairs (i, j) with $i < j$ without duplication
- Naive task decomposition (dividing range $[0, n)$) would create $O(n^2)$ pair computations
- Current implementation partitions the range $[i_0, i_1)$ and checks each task's range against *all* other balls: $\sum_{i=i_0}^{i_1} (n - i - 1)$ comparisons per task
- Actual complexity per task is $O((i_1 - i_0) \cdot n)$, not $O((i_1 - i_0)^2)$

- **Memory Bandwidth Saturation:**

- Task-based approach requires all workers to read shared `Ball` objects' position/velocity data
- With k cores simultaneously accessing the same memory region, cache line contention increases
- `ConcurrentHashMap` operations on `lastTouchedBy` introduce additional memory traffic

- **Synchronization Overhead in Collision Detection:**

```

1   synchronized (first) {
2       synchronized (second) {
3           // Atomic collision resolution
4           Ball.resolveCollision(a, b);
5       }
6   }

```

Listing 1: Ordered Lock Acquisition

- Each collision check acquires locks on two `Ball` objects using `synchronized(first)` then `synchronized(second)`
- Lock ordering (by `System.identityHashCode()`) prevents deadlock but adds comparison overhead
- If two tasks try to resolve the same collision pair with different order, one will block the other

1.3.2 Temporal Constraints

Task-based parallelism introduces timing-related constraints that differ from the thread-per-entity model:

- **Barrier Synchronization Latency:**

- In `WorkerPool`, `PlayerCollisionTask`, `parallelBallUpdate()` and `parallelCollisionDe` use `invokeAll()`, which is a synchronization barrier
- All tasks must complete before control returns to the `GameLoop`
- Frame time = $\max(\text{task duration}) + \text{overhead}$

- **Task Queue Delay:**

- Created tasks are submitted to the `ExecutorService` queue
- If all workers are busy, new tasks wait in queue, delaying their execution
- With a fixed thread pool of size k , tasks may experience queueing delay proportional to queue depth
- This adds non-deterministic latency to each game frame

- **`Future.get()` Blocking Behavior:**

- The main `GameLoop` thread calls `f.get()` (code in class `WorkerPool`) on all futures from `invokeAll()`
- This is a blocking operation—the `GameLoop` cannot proceed until *all* tasks complete
- Any uncaught exception in a task will be re-thrown at `get()`, potentially crashing the loop

1.3.3 Synchronization Requirements

Task-based parallelism requires careful synchronization to ensure memory consistency and avoid data races:

- **`ConcurrentHashMap` for `lastTouchedBy`:**

- Multiple `CollisionTask` instances concurrently update `lastTouchedBy` map via `put()` calls
- Using a non-concurrent map would cause data races
- `ConcurrentHashMap` uses bucket-level locking (segment locking), reducing contention vs. a single global lock
- However, repeated `put()` calls on the same `Ball` key still incur synchronization overhead

- **Read-Write Visibility:**

- After `invokeAll().get()`, all field writes in tasks are visible to the main `GameLoop` (happens-before guarantee)

- The Java Memory Model ensures that synchronization actions (lock acquire/release) create memory barriers
- Without these barriers, updates to `Ball.vel` and `Ball.pos` made by one task might not be visible to the next phase
- **Task Isolation:**
 - `BallUpdateTask` accesses `board` (the `Board` object) to query boundary constraints
 - This creates a shared dependency: if `Board` is modified during task execution, physics updates may read stale boundary data
 - Current implementation assumes `Board` state (bounds, player positions) is immutable during a phase
 - If asynchronous modifications occur (e.g., player ball movement), visibility must be explicitly synchronized
- **Exception Handling and Task Failure:**
 - If a task throws an unchecked exception, `f.get()` will re-throw it via `ExecutionException`
 - The `GameLoop` must handle `InterruptedException` (task thread interrupted) and `ExecutionException` (task computation failed)
 - Without proper handling, the entire game loop may crash
 - Current implementation in `GameLoop.run()` catches both exceptions and breaks the loop gracefully

2 Design and Architecture

2.1 Architectural Overview for Multithreading

The project is structured following the **Model-View-Controller (MVC)** pattern, with a strict separation of responsibilities across three packages: `pcd.assignment01.model`, `pcd.assignment01.view`, and `pcd.assignment01.controller`.

- **Model:** The model package contains all domain state and physics logic, with no knowledge of rendering or threading policy. The central class is `Board`, which owns the list of `Ball` objects, the two player balls, the board boundaries, the scores, and the game-over flag. `updatePlayers()`, `resolveBallCollision()`, `resolvePlayerCollisions()`, and `checkHolesAndGameOver()` are the only methods through which the controller modifies the model state.
- **View:** The view package contains `View`, `ViewFrame`, and `ViewModel`. The view never reads directly from `Board`: it reads exclusively from `ViewModel`, a thread-safe snapshot object whose fields are updated by the controller once per tick via `update(Board, int)`. This decoupling ensures that the Swing Event Dispatch Thread (EDT) always paints a consistent, fully-populated snapshot without ever competing with the physics threads for `Board`'s intrinsic lock.

ViewModel stores **BallViewInfo** records (position and radius) for all small balls and for each player, the current scores, the FPS counter, and the game-over message. The game-over message is written only once, through **setGameOverMessage(String)**, called by the controller at the end of the game; it is never written during normal tick updates, avoiding a race between physics threads and the EDT.

RenderSynch provides a generation counter that lets the update thread block until the EDT has completed the current repaint, preventing frame tearing without busy-waiting.

- **Controller:** The controller package contains **GameController**, **BotController**, **InputHandler**, and **WorkerThread**.

GameController is the orchestrator: it owns the update loop thread, the worker pool, both **CyclicBarrier** instances, and the **BotController**. It is the only component that calls physics methods on **Board** in the correct sequential order after the parallel phases have completed.

BotController encapsulates the autonomous player (player 2) logic in its own thread, waiting on a private monitor until the game physics signal that the ball is still enough to act.

InputHandler translates keyboard events from the EDT into **board.kickPlayer1()** calls, with an early-exit guard on **isGameOver()**. It is registered on the **View** inside **SwingUtilities.invokeLater**, ensuring registration happens on the EDT.




Figure 1: UML with MVC graphics subdivision

2.1.1 Thread Model

The application runs five categories of threads simultaneously during a game:

1. **Main thread:** bootstraps the model and schedules the Swing initialization on the EDT via `SwingUtilities.invokeLater`. It exits immediately after.
2. **Event Dispatch Thread (EDT):** It handles key events (forwarded by the Input Handler) and executes repaints triggered by `panel.repaint()`. It never touches `Board` directly: key events call `board.kickPlayer1()`, which is `synchronized`, and repaints read only from `ViewModel`.
3. **UpdateThread:** a single daemon thread owned by `GameController`. It drives the game loop: it advances the physics pipeline each tick, signals `BotController`, copies state into `ViewModel`, and calls `view.render()`. It also participates in both `CyclicBarriers` as the $(N + 1)$ -th party, allowing it to perform the sequential post-barrier phases without additional synchronization.
4. **WorkerThread-0 ... WorkerThread-N:** a pool that contains exactly $N = \text{availableProcessors}() + 1$ persistent daemon threads, created once at startup and never recreated. Each worker waits on its private `tickMonitor` until `GameController` calls `startTick()`, then executes phase 1 (positions) and phase 2 (ball-ball collisions) before waiting again for the next tick.
5. **BotThread:** a single daemon thread owned by `BotController`. It waits on a private `botMonitor` until `GameController` signals that player 2 is still, then applies a random kick impulse subject to a 2-second cooldown.

All threads except the main thread are daemon threads, so the JVM exits cleanly when the window is closed.

2.1.2 Parallel Physics Pipeline

The physics update for each game tick is divided into five phases, two of which execute in parallel across the worker pool.

1. **Player update (sequential):** Before the parallel work begins, `GameController` calls `board.updatePlayers(dt)`, which advances the positions of the two heavy player balls. This is sequential because (a) there are only two such balls, and (b) the workers' collision phase reads player state indirectly via `resolvePlayerCollisions()`, which runs after the barriers. Keeping player updates sequential avoids any contention.
2. **Position update (parallel):** The list of small balls is divided into N contiguous, disjoint partitions using a balanced split (each partition differs in size by at most one element). Each `WorkerThread` receives its partition and calls `ball.updateState(dt)` on every ball it owns. Because the partitions are disjoint, no two workers write to the same `Ball` object during this phase, and no locking is required.

After finishing, each worker calls `barrierPositions.await()`. The `UpdateThread` also calls `barrierPositions.await()` as the $(N + 1)$ -th party. The barrier scatters all threads into phase 3 only once every worker (and the controller) has finished phase 2, guaranteeing that all positions are fully updated before any collision is resolved.

3. **Spatial grid construction (Worker-0 only):** After all positions have been updated in phase 2, computing ball-ball collisions naively requires checking all $O(n^2)$ pairs. To reduce this to $O(n)$ on average, the application uses a **spatial hashing grid** implemented in the `SpatialGrid` class (package `pcd.assignment01.util`).

The field is partitioned into square cells of side $cellSize = 2 \cdot r_{ball} + \varepsilon$, where $r_{ball} = 0.01$ is the radius of the small balls. This choice guarantees that two balls whose centres lie in non-adjacent cells have distance $> cellSize > 2 \cdot r_{ball}$ and therefore cannot overlap: no colliding pair is ever missed by checking only the 3×3 neighbourhood of cells around each ball.

The grid is stored as a `HashMap<CellKey, List<Ball>>`, where `CellKey` is a value type wrapping the integer column and row indices of a cell. Using a hash map rather than a dense 2D array avoids allocating memory for the large number of empty cells that arise with sparse ball distributions.

The grid must be rebuilt every tick because ball positions change in phase 2. This work is assigned exclusively to **Worker-0** (`isGridBuilder == true`), which calls `grid.build(allBalls)` after `barrierPositions` has fired.

The synchronization for this phase uses a `ReusableLatch` rather than a `CyclicBarrier`. The distinction is deliberate: a `CyclicBarrier` models *symmetric* synchronization — N threads performing parallel work and waiting for each other. Grid construction is instead an *asymmetric* one-to-all dependency: one producer (Worker-0) generates a result that the other workers must wait for before entering phase 4.

Worker-0 calls `gridLatch.open()` after `grid.build()` completes; all other workers call `gridLatch.await()` and block until the latch opens. The `UpdateThread` does not participate in the latch — it proceeds directly from `barrierPositions` to `barrierCollisions`. The `GameController` calls `gridLatch.reset()` at the start of each tick, before activating any worker, making the latch safely reusable across ticks.

The `open()` call establishes a *happens-before* relationship between Worker-0's writes to the grid and all subsequent reads by the other workers in phase 4, as required by the Java Memory Model.

The synchronization points per tick are therefore:

$$\underbrace{\text{barrierPositions}}_{\text{symmetric}} \rightarrow \underbrace{\text{gridLatch}}_{\text{asymmetric}} \rightarrow \underbrace{\text{barrierCollisions}}_{\text{symmetric}}$$

The two `CyclicBarrier` instances are each constructed with $parties = N + 1$ (one per worker plus the `UpdateThread`). The `ReusableLatch` has no party count — it is opened by exactly one thread and awaited by all others independently.

4. **Ball-ball collision resolution with spatial hashing (parallel)**

Each worker iterates over its partition and, for each ball a , calls `grid.getCandidates(a)` to obtain the balls in the 3×3 cell neighbourhood. This returns on average only a small constant number of candidates (proportional to local ball density) instead of all n balls, reducing the total work across all workers from $O(n^2)$ to $O(n)$ on average.

Each worker also builds a `Map<Ball, Integer>` index (populated in `startTick()`) that maps every ball to its position in the shared `allBalls` snapshot. This allows $O(1)$ index lookup during phase 4, replacing the earlier $O(n)$ `indexOf()` call that would have negated the benefit of the spatial grid.

The canonical lock-ordering protocol is preserved: a worker processes the pair (a, b) only if $idx(a) < idx(b)$, ensuring each pair is handled exactly once across all workers. Locks are acquired in ascending index order (`synchronized(a)` then `synchronized(b)`), which eliminates the possibility of deadlock by the same argument as before.

The performance improvement is significant with dense configurations. With the `MassiveBoardConf` (4500 balls), the naive $O(n^2)$ approach requires checking approximately 10^7 pairs per tick. With the spatial grid, each ball checks on average fewer than 10 neighbours, reducing the total to roughly 45 000 candidate pairs — a reduction of more than two orders of magnitude.

5. **Player collision resolution (sequential):** `GameController` calls the method `board.resolvePlayerCollisions()`, which is `synchronized` on `Board`.

This method resolves collisions between each player and every small ball, and between the two players themselves. It also updates the `lastTouchedBy` map, which records the last player to touch each small ball and is used for scoring. This phase is sequential because both player objects are shared across all potential pairs, making fine-grained locking impractical.

6. **Hole detection and game-over (sequential):** `GameController` calls `board.checkHolesAndGameOver()`, which is also `synchronized` on `Board`. It checks whether any ball or player has entered a hole, removes sunk balls from the list, increments scores using `lastTouchedBy`, and sets `isGameOver` if appropriate. This phase is sequential because it mutates the shared `balls` list and the scoring state atomically.

2.1.3 Synchronization Mechanisms

- **CyclicBarrier**

The `CyclicBarrier` (in package `pcd.assignment01.util`) is implemented from scratch using Java's built-in monitor protocol (`synchronized`, `wait()`, `notifyAll()`). It maintains three fields: `parties` (the total number of threads that must call `await()` before the barrier fires), `count` (the number of threads still outstanding), and `generation` (an integer that increments on each firing and is used as the condition variable to distinguish the current cycle from the next one, guarding against spurious wakeups).

When the last thread arrives (`count == 0`), it resets `count` to `parties`, increments `generation`, and calls `notifyAll()`. All waiting threads check `myGeneration == generation` as their loop condition: once the generation advances, they know the barrier has fired and they exit the wait.

The barrier also maintains a `broken` flag. If any waiting thread is interrupted, it sets `broken = true` and calls `notifyAll()`, which causes every other waiting thread to wake up and throw `InterruptedException`. This prevents the pathological scenario where a single interrupted thread leaves `count` permanently non-zero,

causing all other threads to block forever. The `broken` flag is cleared when the last thread successfully fires the barrier for the next generation.

Two instances of `CyclicBarrier` are created per game, each with $parties = N + 1$ (one per worker plus the `UpdateThread`): `barrierPositions` separates phase 2 from phase 3, and `barrierCollisions` separates phase 4 from the sequential phases 5 and 6.

- **ReusableLatch**

The `ReusableLatch` (in package `pcd.assignment01.util`) models *asymmetric* synchronization: one producer thread calls `open()` to signal that a resource is ready, and any number of consumer threads call `await()` and block until the latch is open.

It is implemented with a single boolean flag `open` and the standard monitor protocol. `await()` loops on `while (!open) { wait(); }`, and `open()` sets the flag and calls `notifyAll()`. The `reset()` method closes the latch again for the next cycle; it is called by the `GameController` at the start of each tick, before any worker is activated, so no thread can be in `await()` at reset time.

One instance of `ReusableLatch` is created per game: `gridLatch` gates the transition from phase 3 to phase 4. Worker-0 calls `gridLatch.open()` after `grid.build()` completes; all other workers call `gridLatch.await()`. The `UpdateThread` does not participate in the latch — it proceeds directly from `barrierPositions` to `barrierCollisions`, since it has no work to do during phase 5.

The choice of `ReusableLatch` over a third `CyclicBarrier` is deliberate: a `CyclicBarrier` implies that all parties perform parallel work and wait for each other, which would misrepresent the asymmetric nature of grid construction. Using the semantically correct primitive makes the synchronization structure of the system explicit and self-documenting.

- **Per-ball locking and deadlock prevention:** During phase 4, each worker acquires `synchronized` locks on two `Ball` objects before calling `Board.resolveBallCollision()`. To prevent deadlock, locks are always acquired in ascending index order (the index of the ball in the shared `allBalls` snapshot). Because this order is a strict total order and is respected by all workers, no circular wait can occur: a thread holding `lock(ball[i])` and waiting for `lock(ball[j])` with $i < j$ will never be blocked by another thread holding `lock(ball[j])` and waiting for `lock(ball[i])`.
- **ConcurrentHashMap for lastTouchedBy:** The `lastTouchedBy` map is accessed from two different contexts: from `WorkerThreads` during phase 4 (to remove entries for colliding balls) and from the `UpdateThread` during phases 5 and 6 under the `Board` lock (to update and query entries). Using a `ConcurrentHashMap` makes individual operations on the map thread-safe without requiring a `synchronized(this)` block inside `resolveBallCollision()`. This eliminates a potential deadlock that would arise from the lock ordering `synchronized(ball) → synchronized(board)`, which could interleave incorrectly with `synchronized(board)` held by `resolvePlayerCollisions()`.
- **BotController monitor:** The `BotController` uses a private `Object` `botMonitor` as a condition variable. The `BotThread` waits on this monitor inside a guarded

loop that checks both `running` and `board.isPlayer2Still()`. The `UpdateThread` calls `botController.signal()` once per tick; if player 4 is still, `signal()` calls `notifyAll()` on the monitor. The `stop()` method calls `notifyAll()` unconditionally, ensuring the `BotThread` is always unblocked on shutdown regardless of the ball's current velocity.

- **ViewModel snapshot and thread safety:** All public methods of `ViewModel` are synchronized on the `ViewModel` instance. The `UpdateThread` calls `update()` once per tick, which replaces the internal collections with fresh copies of the model state. The EDT calls getters (`getBalls()`, `getPlayer1()`, etc.) during repaints. Because `update()` returns deep copies (e.g. `new ArrayList<>(balls)`) and all getters return copies as well, the EDT never holds a reference to an object being mutated by the physics threads.

2.1.4 MVC and Multithreading Interaction

The interaction between the MVC layers in the presence of multiple threads is governed by three invariants:

1. **The Model is only mutated by the controller.** The EDT never calls any mutating method on `Board` except `kickPlayer1()`, which is **synchronized**. The `BotThread` calls only `kickPlayer2()`, also **synchronized**. All physics mutations happen exclusively within the sequential and parallel phases orchestrated by `GameController`.
2. **The View never reads the Model directly.** `ViewFrame` and its inner `VisualiserPanel` read exclusively from `ViewModel`. The snapshot is written once per tick by the `UpdateThread` after all physics phases are complete, so the view always paints a consistent state that reflects a completed tick.
3. **Game-over handling is atomic with respect to rendering.** When `checkHolesAndGameOver()` sets `isGameOver = true`, the `UpdateThread` detects this before calling `viewModel.update()` or `view.render()`.

It calls `handleGameOver()`, which writes the winner message to `ViewModel` via `setGameOverMessage()`, and then renders one final frame. This ensures the game-over overlay is shown on a frame that already reflects the final physical state, with no intermediate partially-updated frame visible to the player.

Figure 2 illustrates the interactions between the active threads of the system during a single game tick. Five participants are represented: the `UpdateThread`, which orchestrates the entire tick; `WorkerThread-0`, which additionally acts as the spatial grid builder; `WorkerThread-1..N`, representing all remaining worker threads which behave identically; the `BotThread`, which operates independently; and `Board`, included as a passive participant to make explicit which methods are invoked sequentially and by which thread. The Event Dispatch Thread is intentionally omitted, as it interacts with the system exclusively through **synchronized** methods on `Board` and `ViewModel` and does not participate in any of the barrier synchronization points shown. The diagram covers five phases. In **Phase 1**, the `UpdateThread` sequentially calls `board.updatePlayers(dt)` before any parallel work begins. In **Phase 2**, the `UpdateThread` dispatches `startTick()` to all workers and immediately blocks on `barrierPositions.await()`, while `WorkerThread-0` and




Figure 2: Sequence Diagram of thread interactions and synchronization points during a single game tick.

WorkerThread-1..N update the positions of their respective ball partitions in parallel. The first red dashed line marks the firing of `barrierPositions`, after which all threads are guaranteed to have completed position updates. In **Phase 3**, only **WorkerThread-0** performs work — rebuilding the **SpatialGrid** from the updated positions — while all other threads block immediately on `barrierGrid.await()`. The asymmetry in activation box lengths communicates this visually: **WorkerThread-0** has a tall box while the others have near-zero boxes. The second red dashed line marks the firing of `barrierGrid`, establishing the happens-before relationship that allows all workers to safely read the grid in the next phase. In **Phase 4**, all workers resolve ball-ball collisions in parallel using the spatial grid, acquiring per-ball locks in canonical index order to prevent deadlock, while the **UpdateThread** again blocks immediately. The third red dashed line marks the firing of `barrierCollisions`. Finally, in **Phases 5 and 6**, the **UpdateThread** exclusively calls `board.resolvePlayerCollisions()` and `board.checkHolesAndGameOver()` sequentially, then signals the **BotThread** via `signal()` before updating the **ViewModel** and triggering a render.

2.2 Architectural Overview for Task-Based Approach

The task-based concurrent implementation differs fundamentally from the thread-per-entity model. Rather than creating persistent threads for each game component, this approach uses the **Executor Framework** with a fixed-size thread pool to parallelize computation-heavy phases of the game loop. We used **Map-Reduce** algorithms, the

architecture implements a Master/Worker pattern where the **GameLoop** acts as the master coordinator. It delegates tasks to a **WorkerPool** abstraction which, knowing the data space topology in advance, splits the workload into smaller, independent chunks (*map phase*). These chunks are then efficiently computed in parallel by the underlying **ExecutorService**.




Figure 3: UML class diagram

2.2.1 Component Descriptions

The task-based architecture consists of the following key components:

GameLoop (Master/Orchestrator Thread)

The **GameLoop** extends **Thread** and serves as the main control flow:

- **Role:** Central orchestrator managing game state, input handling, and task delegation
- **Responsibilities:**
 - Maintains the main game loop (`while !board.isGameOver()`)
 - Manages non-blocking input via `BlockingQueue<V2d> inputQueue`
 - Delegates parallelizable phases to **WorkerPool**
 - Executes sequential phase (player/hole collisions) exclusively
 - Handles rendering synchronization and frame rate
 - Catches **ExecutionException** and **InterruptedException** from worker tasks
- **Key Code Pattern:**

```

1  while (!board.isGameOver()) {
2      V2d impulse = inputQueue.poll(); // Non-blocking input
3
4      // Phase 1: Parallel physics
5      workerPool.parallelBallUpdate(board.getBalls(), elapsed,
6                                     board);
7
8      // Sequential player logic
9      if (board.getPlayer1() != null)
10         board.getPlayer1().updateState(elapsed, board);
11
12     // Phase 2: Parallel collision detection
13     workerPool.parallelCollisionDetection(board.getBalls(),
14                                           board.getLastTouchedBy());
15
16     // Phase 3: Parallel player updates
17     workerPool.parallelPlayerCollision();
18
19     board.checkHolesAndEndGame();
20
21     // Render
22     viewModel.update(board, fps);
23     view.render();
24 }

```

Listing 2: GameLoop Main Loop Structure

WorkerPool (Task Coordinator Abstraction)

The WorkerPool encapsulates all ExecutorService framework details from the GameLoop:

- **Role:** Hides complexity of task scheduling, execution, and synchronization
- **Internal Structure:**
 - ExecutorService executor: Fixed thread pool with size $k = 2 \times \text{availableProcessors}$
 - Three public methods: `parallelBallUpdate()`, `parallelCollisionDetection()`, and `parallelPlayerCollision()`
 - All use `invokeAll()` for barrier synchronization
- **Key Guarantee:** The GameLoop does not import `java.util.concurrent.*`, maintaining clean separation of concerns
- **Exception Propagation:** Exceptions in tasks are bundled into `ExecutionException` and re-thrown at `f.get()`

BallUpdateTask (Parallel Worker)

Implements `Callable<Void>` to update a range of balls' physics state:

- **Input Parameters:**

- `List<Ball> balls`: Shared reference to all balls
- `int from, int to`: Range $[from, to)$ of balls this task processes
- `long dt`: Elapsed time since last frame
- `Board board`: Context for boundary constraints

- **Computation:**

```

1 public Void call() {
2     for (int i = from; i < to; i++) {
3         balls.get(i).updateState(dt, board);
4     }
5     return null;
6 }

```

Listing 3: BallUpdateTask Execution

- **Synchronization:** No locks required; each task reads from distinct list indices and modifies only the Ball objects at those indices. The Board's boundary data is assumed immutable during execution.

CollisionTask (Parallel Worker)

Implements `Callable<Void>` to resolve pairwise ball collisions:

- **Input Parameters:**

- `List<Ball> balls`: Shared reference to all balls
- `int from, int to`: Range $[from, to)$ of "first" balls in pairs
- `ConcurrentHashMap<Ball, Integer> lastTouchedBy`: Shared map of ball-to-player touches

- **Pair Enumeration:** For each ball $i \in [from, to)$, checks pairs (i, j) where $j \in [i + 1, n)$

- **Synchronization Protocol:**

```

1 for (int i = from; i < to; i++) {
2     for (int j = i + 1; j < n; j++) {
3         Ball a = balls.get(i);
4         Ball b = balls.get(j);
5
6         // Ordered lock acquisition (deadlock prevention)
7         Ball first = System.identityHashCode(a) <= System.
            identityHashCode(b) ? a : b;
8         Ball second = (first == a) ? b : a;
9
10        synchronized (first) {
11            synchronized (second) {
12                Ball.resolveCollision(a, b);
13                // Update lastTouchedBy if collision occurred
14                if (velocities changed) {

```

```

15         lastTouchedBy.put(a, 0);
16         lastTouchedBy.put(b, 0);
17     }
18 }
19 }
20 }
21 }

```

Listing 4: Ordered Lock Acquisition in CollisionTask

- **Fine-Grained Locking:** Each pair's collision resolution is protected by locks on both Ball objects, but independent pairs may execute concurrently.

PlayerCollisionTask (Parallel Worker)

Implements `Callable<Void>` to resolve player-ball collisions in parallel:

- **Input Parameters:**
 - `List<Ball> balls`: Shared reference to all balls
 - `int from, int to`: Range $[from, to)$ of "first" balls in pairs
 - `ConcurrentHashMap<Ball, Integer> lastTouchedBy`: Shared map of ball-to-player touches
- **Pair Enumeration:** For each ball $i \in [from, to)$, checks pairs (i, j) where $j \in [i + 1, n)$
- **Synchronization Protocol:**

```

1  for (int i = from; i < to; i++) {
2      Ball b = balls.get(i);
3
4      // Collisione con player1
5      if (player1 != null) {
6          V2d vPrima = b.getVel();
7          Ball.resolveCollision(player1, b);
8          if (!b.getVel().equals(vPrima)) {
9              lastTouchedBy.put(b, 1);
10         }
11     }
12
13     // Collisione con player2
14     if (player2 != null) {
15         V2d vPrima = b.getVel();
16         Ball.resolveCollision(player2, b);
17         if (!b.getVel().equals(vPrima)) {
18             lastTouchedBy.put(b, 2);
19         }
20     }
21 }
22 return null;

```

Listing 5: Ordered Lock Acquisition in PlayerCollisionTask

Task Decomposition

Task decomposition in the Pool implementation follows a **range-partitioning** strategy:

- **Chunk Size Calculation:**

```

1  int nWorkers = Runtime.getRuntime().availableProcessors() *
    2;
2  int chunkSize = Math.max(1, n / nWorkers);

```

Listing 6: Chunk Size Computation in WorkerPool

where n is the number of balls. This ensures:

- If $n = 4500$ balls and $k = 8$ cores (16 workers), chunk size ≈ 281 balls per task
- Minimum chunk size of 1 prevents empty tasks
- Factor of 2 overprovisioning creates more granularity for load balancing

2.3 Synchronization Strategy

Task-based synchronization employs a layered strategy combining barrier synchronization with fine-grained locking:

2.3.1 Barrier Synchronization via invokeAll()

Parallel phases use `ExecutorService.invokeAll()`, which provides implicit barrier semantics:

- **Behavior:** Blocks until *all* submitted tasks complete
- **Java Memory Model Guarantee:** All field writes in tasks are visible after `invokeAll()` returns (happens-before relationship)
- **Game Loop Implications:**
 - Phase 1 completes $\implies$ all ball positions/velocities updated and visible
 - Phase 2 completes $\implies$ all collision resolutions applied and visible
 - Phase 3 completes $\implies$ all player updates applied and visible
 - Sequential Phase 4 executes with consistent global state
- **Performance Cost:** Tail latency — the slowest task determines phase duration

2.3.2 Fine-Grained Ball Locking in Collision Detection

`CollisionTask` uses nested synchronized blocks to prevent data races during collision resolution:

- **Two Locks:** Each `Ball.resolveCollision(a, b)` modifies both `Ball` objects' velocity and position. Without synchronization, a concurrent update might corrupt the ball state.
- **Deadlock Prevention via Ordering:**

- All code acquires locks in a consistent order: always lock Ball A before Ball B if `identityHashCode(A) ≤ identityHashCode(B)`
- Since identity hash codes are fixed per object, this prevents circular wait (requirement for deadlock)
- **Lock Contention:** Multiple tasks may compete for the same Ball locks
 - If two tasks both need to resolve collisions involving the same Ball, one task waits
 - This serializes independent collisions unnecessarily, but ensures safety
 - In practice, contention is moderate since each task processes a different range of "first" balls

2.3.3 ConcurrentHashMap for Shared State

The `lastTouchedBy` map is a `ConcurrentHashMap<Ball, Integer>` to allow concurrent updates from multiple `CollisionTask` instances:

- **Concurrent Put Operations:** Multiple tasks call `lastTouchedBy.put(ball, 0)` simultaneously
- **Bucket-Level Locking:** `ConcurrentHashMap` uses segment-based locking (in modern Java, striped locking)
 - Reduces contention compared to a single global synchronized lock
 - Different Ball keys typically hash to different buckets, allowing concurrent puts
 - Same-bucket updates still serialize
- **Visibility Guarantee:** All puts are visible globally via `ConcurrentHashMap`'s internal synchronization

2.3.4 Exception Handling and Reliability

Task-based execution requires careful exception handling to prevent loop crashes:

- **ExecutionException:** If a task throws an unchecked exception, `f.get()` re-throws it wrapped in `ExecutionException`
- **InterruptedException:** If the `GameLoop` thread is interrupted (e.g., window closed), waiting on futures throws `InterruptedException`
- **GameLoop Protection:**

```

1  try {
2      workerPool.parallelBallUpdate(board.getBalls(), elapsed,
          board);
3      workerPool.parallelCollisionDetection(board.getBalls(),
          board.getLastTouchedBy());
4      workerPool.parallelPlayerCollision();
5  } catch (InterruptedException | ExecutionException e) {
6      Thread.currentThread().interrupt();
7  }

```

```

8      break; // Exit game loop gracefully
9  }

```

Listing 7: Exception Handling in GameLoop

- **Graceful Degradation:** Catching exceptions allows the game loop to terminate cleanly rather than crashing the entire JVM

2.3.5 Memory Visibility and Happens-Before Guarantees

The Java Memory Model ensures correctness through explicit synchronization points:

1. **Task Submission:** When a task is submitted to the executor, all field writes in the GameLoop thread are visible to the task
2. **Task Execution:** Workers read shared Ball objects and Board state; without locks, visibility is not guaranteed
3. **invokeAll() Barrier:** When invokeAll() returns, all task writes are visible to the GameLoop thread
4. **Nested Locks in Collisions:** Each `synchronized` block creates a memory barrier, ensuring Ball field updates are visible across tasks

This ensures that despite the distributed nature of tasks, the game state remains consistent across phase boundaries.

3 Behavior Specification: Petri Nets

3.1 Multi-Threaded CyclicBarrier Net

The `CyclicBarrier` is a custom synchronization primitive implemented from scratch. Modelling it as a Petri Net allows its correctness properties — absence of deadlock, liveness, and automatic reset — to be verified independently of the application logic that uses it.

The net models a barrier with $N = 3$ parties (the number used during JPF verification: two `WorkerThreads` and the `UpdateThread`). Each party is represented by a dedicated input place P_i carrying an initial token, and a dedicated output place P'_i . A single transition *barrier fires* consumes one token from each input place and produces one token in each output place.

The key properties visible from the net structure are the following.

Absence of deadlock. The transition *barrier fires* is enabled if and only if all N input places simultaneously hold a token. Since every party eventually deposits its token (no alternative path exists in the net), the transition is guaranteed to fire. No subset of parties can block the others indefinitely.

Liveness. After *barrier fires*, every output place P'_i receives a token, allowing all parties to proceed. The net is live: every transition that is ever enabled will eventually fire in any fair execution.

Cyclic reset. The output places P'_i are connected back to the input places P_i of the next barrier cycle (not shown in the figure for clarity), restoring the initial marking and making the barrier ready for the next tick without reallocation.

Boundedness. The net is 1-safe: at most one token can reside in any place at any time. This follows from the fact that each party deposits exactly one token per cycle, and the transition consumes all tokens atomically.

3.2 Multi-Threaded Tick Flow

The second net models the complete flow of a single game tick across all active threads: the `UpdateThread`, `WorkerThread-0`, `WorkerThread-1..N`, and the `BotThread`. Each thread is represented as a horizontal row of places connected by transitions; vertical transitions that span multiple rows model synchronization points between threads.




Figure 4: Petri Net modelling the per-thread state flow during a single game tick. Vertical transitions crossing multiple rows represent synchronization points. The `gridLatch` transition spans only the worker rows, reflecting its asymmetric producer-consumer semantics.

Thread states. Each place in the net represents a meaningful state of the corresponding thread. For the `UpdateThread`: *tick start*, *update players*, *await barrierGrid*, *sequential phases*, *rendering*, and *tick done*. For `WorkerThread-0`: *await startTick*, *phase 1*, *grid build*, *phase 2*, and *await next tick*. For `WorkerThread-1..N`: the same sequence except that *grid build* is replaced by a direct arc from `barrierPositions` to `gridLatch`, reflecting the absence of any work during phase G. For the `BotThread`: *await signal* and *kick (if still)*.

Symmetric synchronization — `barrierPositions` and `barrierCollisions`. These two transitions span the rows of all three physics threads (`UpdateThread`, `WorkerThread-0`, and `WorkerThread-1..N`). Each is enabled only when all three rows simultaneously hold a token in the corresponding input place, modelling the all-to-all semantics of the `CyclicBarrier`: every thread must complete its current phase before any thread may advance to the next.

Asymmetric synchronization — `gridLatch`. This transition spans only the two worker rows, not the `UpdateThread` row. Its semantics differ fundamentally from those of the two barriers: it models a one-to-all dependency where `WorkerThread-0` (the producer) opens the latch after `grid.build()` completes, and `WorkerThread-1..N` (the consumers) are unblocked. The `UpdateThread` bypasses the latch entirely, proceeding directly from `barrierPositions` to `barrierCollisions` via its own arc. The asymmetry is visible in the net structure: `WorkerThread-0` has an intermediate place *grid build* before the `gridLatch` transition, while `WorkerThread-1..N` has no such place, reflecting that only Worker-0 performs work during phase G.

BotThread independence. The `BotThread` row is decoupled from the physics rows: it has no shared transitions with the barrier or latch synchronization points. It is connected to the rest of the net only by the dashed `signal()` arc from the `UpdateThread`, which fires after `barrierCollisions` and triggers the bot's kick transition if `isPlayer2Still()` holds. This independence is a deliberate design choice: the bot must not interfere with the physics pipeline timing.

Behavioral properties. The net is 1-safe and live by the same argument as the `CyclicBarrier` net: every thread row has exactly one token in circulation at any time, and every synchronization transition is eventually enabled because no thread has an alternative path that bypasses a barrier or the latch. The loop arcs returning each thread to its initial place after a completed tick confirm that the system can execute indefinitely without reaching a dead marking, consistent with the `no errors detected` result from JPF verification.

3.3 High-Level Game Flow Net Task-based

The system can be broken down into actions called “tasks”; agents are responsible for carrying them out. The computations required to complete each individual task can occur in parallel; however, there are temporal dependencies within the system. The sequential nature of the problem is fundamentally due to the data dependencies between tasks: if a variable must be read in order to complete a task, then it must be ensured that this variable does not change during the read operation, so as not to lose information.

Tasks	Temporal Dependencies	Data Dependencies
Ball Update	Must complete BEFORE Collision Detection	Reads: position, velocity, bounds, elapsed time. Writes: updated position, velocity
Collision Update	Starts AFTER Ball Update barrier; Completes BEFORE Player Collision barrier	Reads: position, velocity, radius, mass (all balls), lastTouchedBy. Writes: velocity, position, lastTouchedBy
Player Collision	Starts AFTER Collision Detection barrier; Completes BEFORE Rendering	Reads: position, velocity, player positions, bounds, lastTouchedBy, scores. Writes: velocity, position, lastTouchedBy, scores, balls list, isGameOver, winnerMessage
Simulation Rendering	Starts AFTER Player Collision; Completes BEFORE next frame input poll	Reads: balls list, player positions, scores, isGameOver, winnerMessage, framePerSec. Writes: ViewModel state, GUI rendering

By following the table, it is possible to represent the basic behavior of the system using a generic Petri net, which is ideally suited for representing the temporal dependencies of individual tasks.




Figure 5: High-Level Game Flow Petri Net

4 Performance Evaluation

4.1 Experimental Setup

4.1.1 Test Configurations

Three board configurations are tested:

Configuration	Small Balls	Expected Complexity
Minimal	2	
Large	400	
Massive	4,500	

Table 1: Board Configurations and Collision Detection Complexity

4.1.2 Performance Metrics

- **Frame Rate (FPS):** Frames rendered per second
- **Update Latency:** Time to compute one physics update cycle
- **CPU Utilization:** Percentage of available cores used
- **Throughput:** Simulation steps per second
- **Speedup:** Concurrent vs. sequential frame rate ratio

4.2 Expected Performance Improvements

4.2.1 Sequential Baseline Issues

The sequential implementation exhibits:

- Minimal/Large: fps (acceptable)
- Massive: 2 fps (unacceptable frame rate degradation)

4.2.2 Concurrent Improvements

Update Thread Parallelization:

- Potential to parallelize ball state updates (independent operations)
- Potential to parallelize collision detection
- Rendering no longer blocks physics computation

Bot Thread Independence:

- Bot decisions computed asynchronously
- Reduces decision latency from main game loop
- Enables responsive user input even during intensive physics

Expected Speedup Formula:

$$S = \frac{T_{seq}}{T_{conc}} = \frac{f_{seq}}{f_{conc}}$$

For massive configuration with parallel collision detection over k cores:

$$S \approx 1 + \frac{(O(n^2) - O(n))}{k}$$

With $n = 4,500$ balls and $k = 4$ cores:

$$S \approx 1 + \frac{10.1M - 4,500}{4} \approx 2.5\text{-}2.8\times$$

// Note: Actual speedup may be lower due to synchronization overhead and non-parallelizable code sections.

4.2.3 Task-based Improvements

5 Formal Verification with JPF

5.1 Java Path Finder Integration

JPF (Java Path Finder) is a model checker developed at NASA for verifying Java programs. The Pool implementation integrates JPF verification through configuration and test drivers.

5.1.1 JPF Configuration

Key aspects of JPF application to Pool:

```
1 target=pcd.assignment01.controller.MainJPF
2 listener = gov.nasa.jpf.listener.PropertyListener
```

Listing 8: JPF Verification Approach

5.2 Safety Properties Verified

5.2.1 1. Mutual Exclusion (No Data Races)

Property: Board state modifications are atomic and do not interleave.

Verification Method:

- JPF tracks synchronized block boundaries
- No two threads can execute synchronized Board methods concurrently
- All field accesses to mutable state (balls, scores) occur within synchronization

Critical Section Analysis:

```
1 public synchronized void updateState(long dt) {
2     // CRITICAL SECTION START
3     for (Ball b : balls) { /* ... */ } // Protected access to
        balls
4     score1++; // Protected access to score
5     // CRITICAL SECTION END
6 }
```

Listing 9: Critical Section: Board.updateState()

5.2.2 2. Deadlock Freedom

Property: No circular wait-for cycles among threads.

Lock Hierarchy:

Lock Order	Acquisition Pattern
1. Board	Always acquired first (main game state)
2. Bot Monitor	Acquired only from within Bot thread
3. RenderSync	Acquired from Update thread only

Table 2: Lock Hierarchy for Deadlock Prevention

Analysis:

- Bot thread never acquires Board after waiting on botMonitor
- No wait chains among Update, Bot, and EDT threads
- Circular dependencies are structurally impossible

5.2.3 3. Liveness (Game Termination)

Property: Game always terminates within bounded time steps.

Termination Conditions:

1. **Score-based:** All small balls scored or removed
2. **Player-based:** A player's ball enters a hole
3. **User-initiated:** Window closed or game stopped

JPF Verification:

```

1 public void testGameTermination() {
2     // JPF will explore all possible execution paths
3     // and verify that all paths eventually reach:
4     assert !isRunning() : "Game must terminate";
5     assert gameOverFlag : "Game-over state must be set";
6 }

```

Listing 10: Termination Assertion

5.3 Race Condition Detection

5.3.1 Synchronized Field Accesses

JPF verifies that all shared mutable fields are protected:

```

1 private List<Ball> balls;           // Protected by synchronized
   methods
2 private int score1;                 // Protected by synchronized
   methods
3 private boolean running;           // Protected by synchronized
   getter/setter

```

```

4 private boolean isGameOver;          // Protected by synchronized
   methods

```

Listing 11: Synchronized Field Protection

5.3.2 Temporal Assertions

Critical invariants verified during execution:

```

1 // Invariant 1: Score consistency
2 assert score1 >= 0 && score2 >= 0 : "Scores are non-negative";
3
4 // Invariant 2: Ball list validity
5 assert balls != null : "Ball list always exists";
6 assert balls.stream().allMatch(b -> b != null) :
7     "No null balls in list";
8
9 // Invariant 3: Game state consistency
10 assert !(isGameOver && !running) :
11     "Game-over implies not running";

```

Listing 12: Invariants Checked by JPF

5.4 Model Checking Results

5.4.1 Test Configuration: MainJPF

The MainJPF class provides a minimal test case for JPF verification:

```

1 public class MainJPF {
2     public static void main(String[] args) {
3         // Minimal board configuration for JPF
4         Board board = new Board();
5         board.init(new MinimalBoardConf());
6
7         GameController controller =
8             new GameController(board, null, null);
9
10        // JPF explores all possible interleavings
11        controller.startGame();
12
13        // Verification terminates after bounded steps
14        try {
15            Thread.sleep(100);
16        } catch (InterruptedException e) {
17            // Expected: JPF may interrupt
18        }
19
20        controller.stopGame();
21    }
22 }

```

Listing 13: MainJPF Verification Driver

References

- [1] NASA Formal Methods Program, *Java Path Finder - Model Checker for Java Programs*, <https://javapathfinder.github.io/>
- [2] J.L. Peterson, *Petri Nets*, *ACM Computing Surveys*, vol. 9, no. 3, 1977.
- [3] C.A.R. Hoare, *Monitors: An Operating System Structuring Concept*, *Communications of the ACM*, 1974.
- [4] B. Goetz et al., *Java Concurrency in Practice*, Addison-Wesley, 2006.